

YVR Energy Use Forecast

Yilun Lu

3/29/2020

Objective:

The main goal of this project is to develop a model to forecast monthly energy use for the Vancouver International Airport (YVR), for the next three years. If possible, we will also be investigating the reasons behind the unexpected surge or drop in energy use. Such might give us actionable insights which might aid airport management in decision making.

Background Information:

Second largest airport in Canada, Vancouver International Airport (YVR) requires huge amount of energy to light up, heat or cool, and control the massive facility each month. The energy cost usually accounts for a great portion of total expenses for airports in general. Therefore, being able to forecast the energy use accurately would better help financial analysts identify any future cost-related financial issues in the airport.

Approaches:

We are going to build a time series model to forecast YVR's energy use for the future months in the next three years. The best model will be selected from the ETS and ARIMA frameworks, based on residual diagnostics, generalization ability on the test set, and model complexity.

Bibliographical:

“U.S. Energy Information Administration - EIA - Independent Statistics and Analysis.” How Much Electricity Does an American Home Use? - FAQ - U.S. Energy Information Administration (EIA), www.eia.gov/tools/faqs/faq.php?id=97&t=3.

```
library(fpp2)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: forecast
```

```
## Registered S3 method overwritten by 'quantmod':
```

```
##   method          from
```

```
##   as.zoo.data.frame zoo
```

```
## Registered S3 methods overwritten by 'forecast':
```

```
##   method          from
```

```
##   fitted.fracdiff  fracdiff
```

```
##   residuals.fracdiff fracdiff
```

```
## Loading required package: fma
```


```
## Loading required package: expsmooth
```

Introduction:

```
df = read.csv("~/Documents/MBAN/BABS502/project/Energy_use_at_YVR.csv")


yvr = ts(df[, 2], start=c(1997, 1), frequency = 12) # only interested in the energy use column

# time plot
plot(yvr, xlab='Years', ylab='Monthly Energy Use at YVR (1000s kWh)', main='Monthly Energy Use at YVR',
```




```
# STL Decomposition
fit.stl <- stl(yvr, t.window=15,
               s.window="periodic", robust=TRUE)
plot(fit.stl, main='STL Decomposed Components')
```

STL Decomposed Components




```
plot(fit.stl$time.series[,1], xlim=c(1998.0, 1999.5))
```



The discussion refers to the STL decomposition plot.

- 1) Cause for the faster increase in trend from 2007 It could be that, more and more people are coming to Vancouver to live, study and do business. Therefore the increases in airlines raise the energy use.
- 2) Cause for the cycle from 1998 to 2003 and the unexpected energy drops between 1999 and 2000 Such is likely to be caused by the creation of Energy Reduction Committee in 1999 to reduce energy use. The creation of committee might also explain the cycle from 1998 to 2003, which indicate the committee's effort was paid-off only temporarily.
- 3) Explanation for the annual seasonality The seasonal pattern is quite intuitive. The energy use starts to rise in spring, reaches its peak in summer, drops again in fall, and eventually rises again in winter. Overall the energy consumption in summer is much greater than that in winter, probably because Vancouver's winter does not need too much heating in the airport.


```
yvr.boxcoxed = BoxCox(yvr, lambda=0.5)
plot(yvr.boxcoxed, xlab='Years', ylab='Boxcoxed Monthly Energy Use at YVR (1000s kWh)', main='Boxcoxed Monthly Energy Use at YVR (1000s kWh)')
```



The purpose of the BoxCox Transformation is to even the seasonal variance across the times series so that we can better model the trend cycle (although in our case the variance isn't that heterogenous). The result of the box cox transformation isn't significant, which indicates that we might not have heterogenous seasonal variance in the first place.

```
monthdays <- rep(c(31,28,31,30,31,30,31,31,30,31,30,31),14)
monthdays[26 + (4*12)*(0:2)] <- 29
yvr.calender.adjusted = yvr/monthdays
```

```
plot(yvr.calender.adjusted, xlab='Years', ylab='Average Dailly Energy Use at YVR (1000s kWh)', main='C
```



The calendar adjustment is to remove the calendar effect in our time series due to variation of number of days in each month. Therefore, by turning the monthly average energy use into daily average we expect to see the time series to be smoothened out. However, what we actually observe is that there seems to be more variations in our time series (graph above). Such means the calendar adjustment is ineffective.

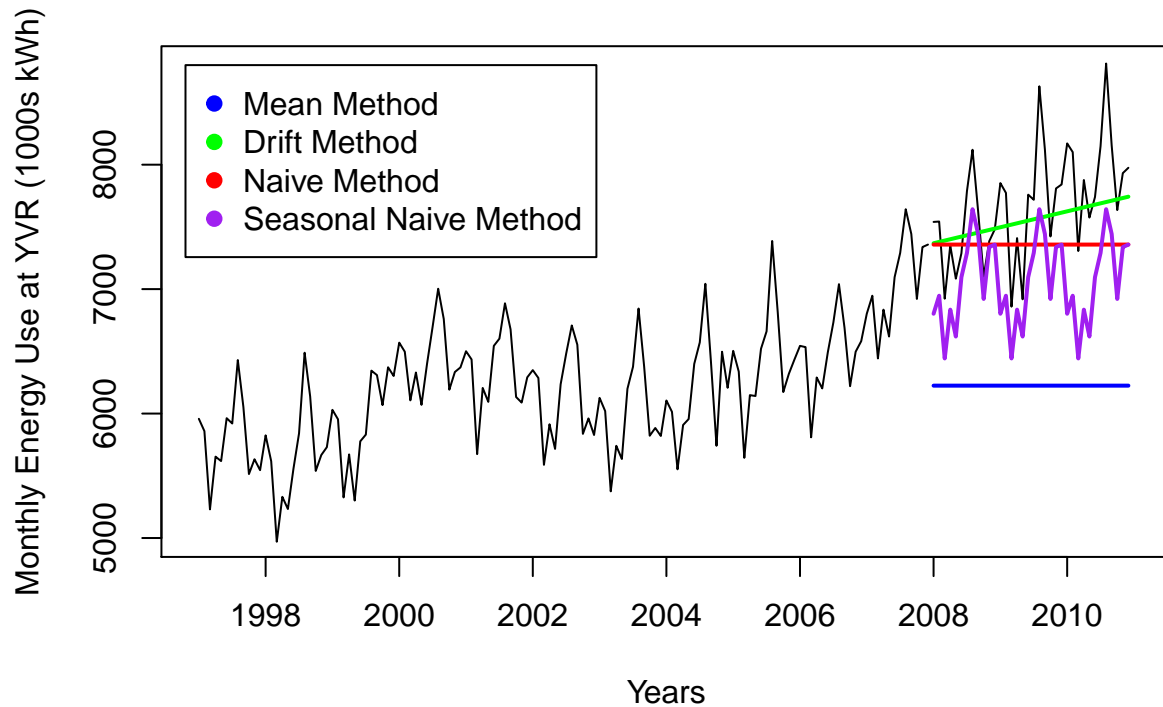
```
yvr.train = window(yvr, end=c(2007, 12))
yvr.test = window(yvr, start=c(2008, 1))
```

```
# mean method (blue)
fit1.mean = meanf(yvr.train, h=36)
# drift method (green)
fit1.drift = rwf(yvr.train, drift = TRUE, h=36)
# naïve method (red)
fit1.naive = naive(yvr.train, h=36)
# seasonal naïve method (purple)
fit1.snaive = snaive(yvr.train, h=36)

# visualize the forecasts
plot(yvr.train, xlim=c(1997, 2011), ylim=c(5000, 8800), main='YVR Test data forecast',
     xlab='Years', ylab='Monthly Energy Use at YVR (1000s kWh)')
lines(yvr.test)
lines(fit1.mean$mean, col='blue', lwd=2)
lines(fit1.drift$mean, col='green', lwd=2)
lines(fit1.naive$mean, col='red', lwd=2)
lines(fit1.snaive$mean, col='purple', lwd=2)
legend(1996.8, 8800, legend=c('Mean Method', 'Drift Method', 'Naive Method', 'Seasonal Naive Method'),
```

```
col=c('blue', 'green', 'red', 'purple'), pch=19)
```

YVR Test data forecast



```
# accuracy measures
```

```
accuracy(meanf(yvr.train, h=36), yvr.test)
```

```
##                ME      RMSE      MAE      MPE      MAPE
## Training set -1.930644e-13  507.840  406.4286 -0.6657438  6.579549
## Test set      1.463795e+03 1528.498 1463.7955 18.7755758 18.775576
##                MASE      ACF1 Theil's U
## Training set  1.533789 0.7140918      NA
## Test set      5.524104 0.4037248  3.138214
```

```
accuracy(naive(yvr.train, h=36), yvr.test)
```

```
##                ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  10.68702 371.8135 310.2901 -0.02398732 5.053622 1.170980
## Test set      330.08333 550.0588 443.1944  3.98109394 5.603453 1.672537
##                ACF1 Theil's U
## Training set -0.1845823      NA
## Test set      0.4037248  1.11998
```

```
accuracy(rwf(yvr.train, h=36, drift = TRUE), yvr.test)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -2.638220e-13 371.6598 309.3927 -0.1967793 5.044069 1.167593
## Test set      1.323734e+02 415.9790 328.4824  1.4437756 4.213286 1.239634
##               ACF1 Theil's U
## Training set -0.1845823      NA
## Test set      0.2664233 0.8509942
```

```
accuracy(snaive(yvr.train, h=36), yvr.test)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 127.9833 343.3462 264.9833 1.853620 4.164394 1.000000
## Test set     626.8333 698.6573 626.8333 8.038624 8.038624 2.365558
##               ACF1 Theil's U
## Training set 0.8206604      NA
## Test set     0.4806562 1.424002
```

The drift method generalizes the best since has the lowest errors on test data, for all metrics.

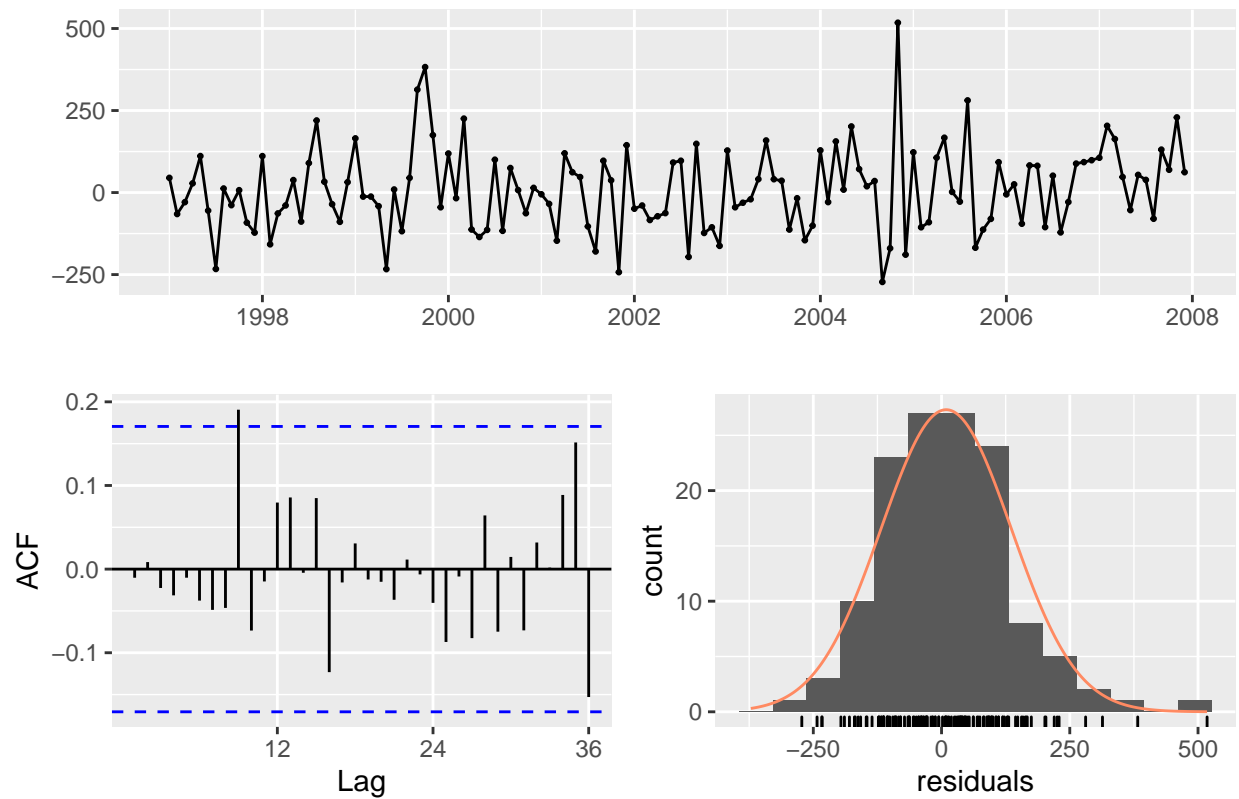
MASE of drift method: on average, the test error, scaled on training set mean absolute error, is 1.2. It means, on average, the test error is 20 percent higher than the training error.

```
fit1.ets = ets(yvr.train, model = 'AAA', damped = TRUE)
fit2.ets = ets(yvr.train, model = 'MAA', damped = TRUE)
fit3.ets = ets(yvr.train, model = 'AAA', damped = NULL)
fit4.ets = ets(yvr.train, model = 'MAA', damped = NULL)
```

```
fit5.ets = ets(yvr.train, model = 'MAM', damped=TRUE)
fit6.ets = ets(yvr.train, model = 'MAM', damped=NULL)
fit7.ets = ets(yvr.train)
```

```
checkresiduals(fit1.ets)
```

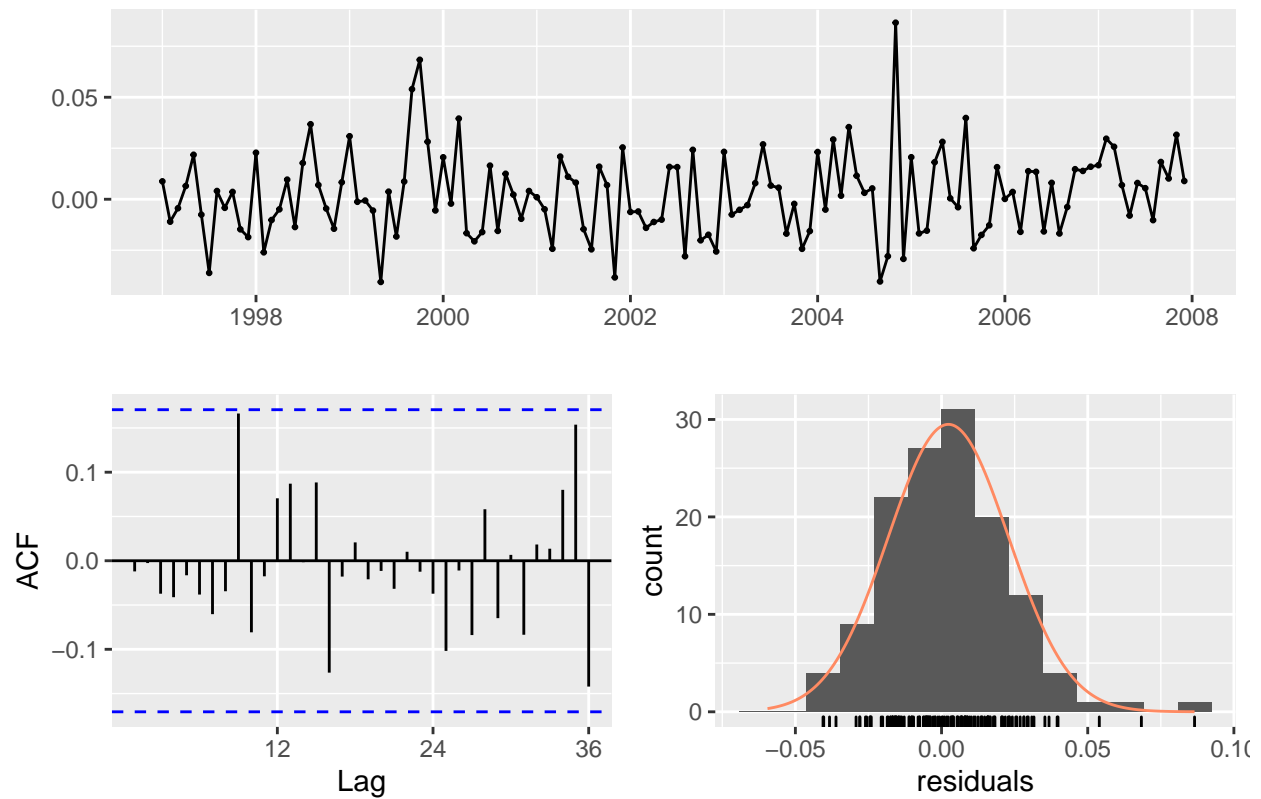

Residuals from ETS(A,Ad,A)



```
##
##  Ljung-Box test
##
## data:  Residuals from ETS(A,Ad,A)
## Q* = 13.308, df = 7, p-value = 0.06495
##
## Model df: 17.    Total lags used: 24
```

```
checkresiduals(fit2.ets)
```

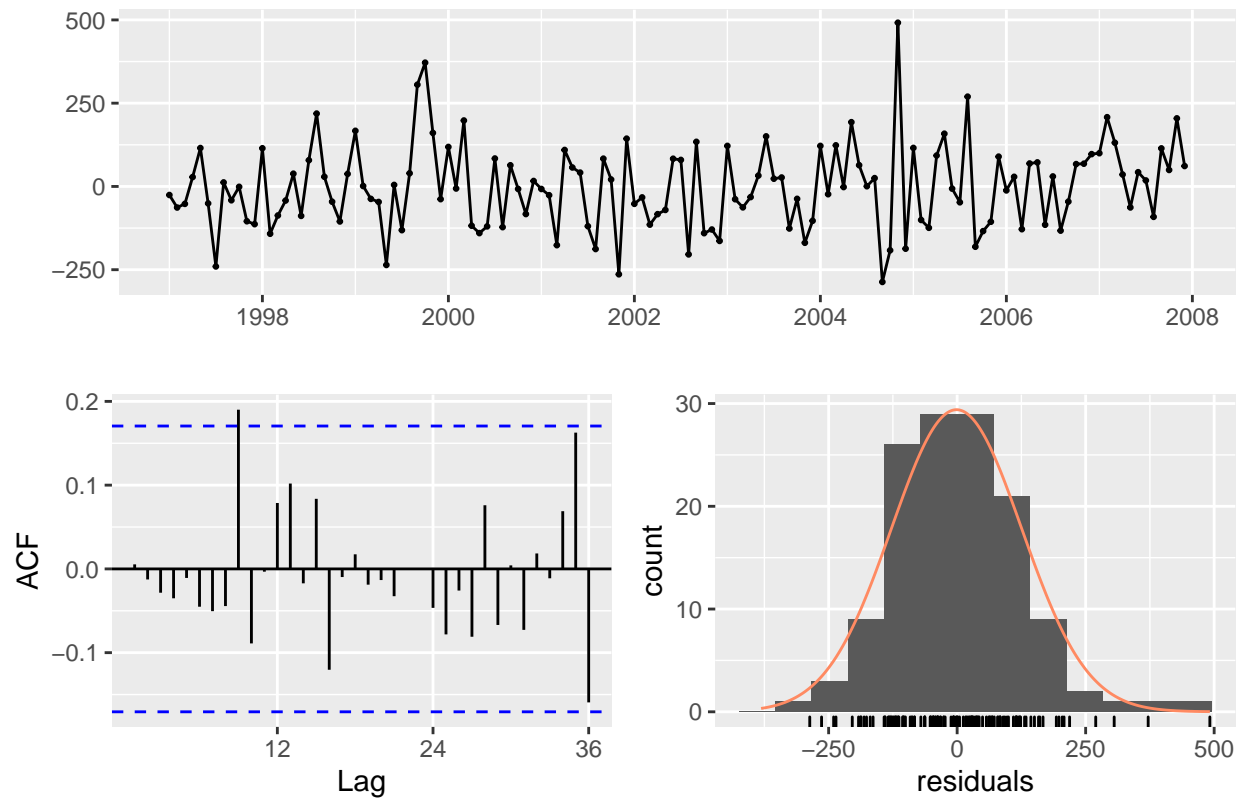
Residuals from ETS(M,Ad,A)



```
##
##  Ljung-Box test
##
## data:  Residuals from ETS(M,Ad,A)
## Q* = 12.451, df = 7, p-value = 0.08667
##
## Model df: 17.   Total lags used: 24
```

```
checkresiduals(fit3.ets)
```

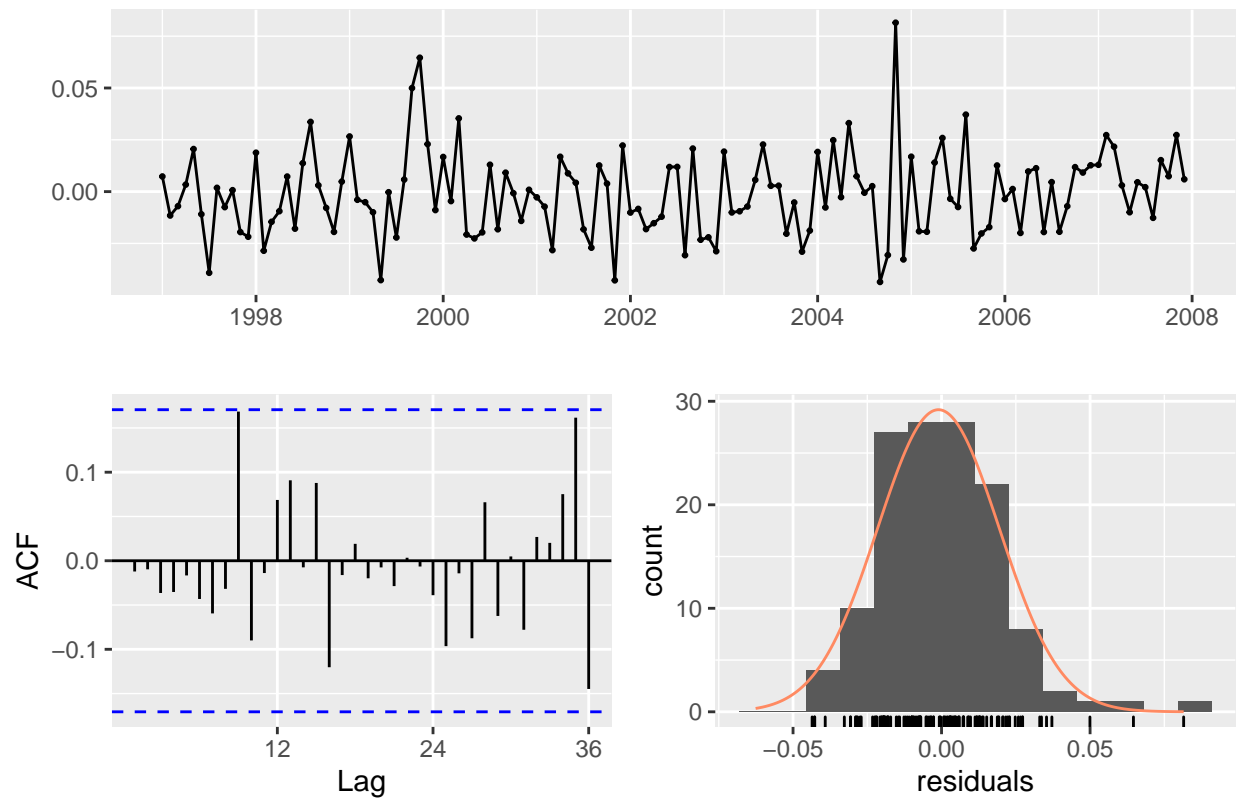
Residuals from ETS(A,A,A)



```
##
##  Ljung-Box test
##
## data:  Residuals from ETS(A,A,A)
## Q* = 14.031, df = 8, p-value = 0.08095
##
## Model df: 16.    Total lags used: 24
```

```
checkresiduals(fit4.ets)
```

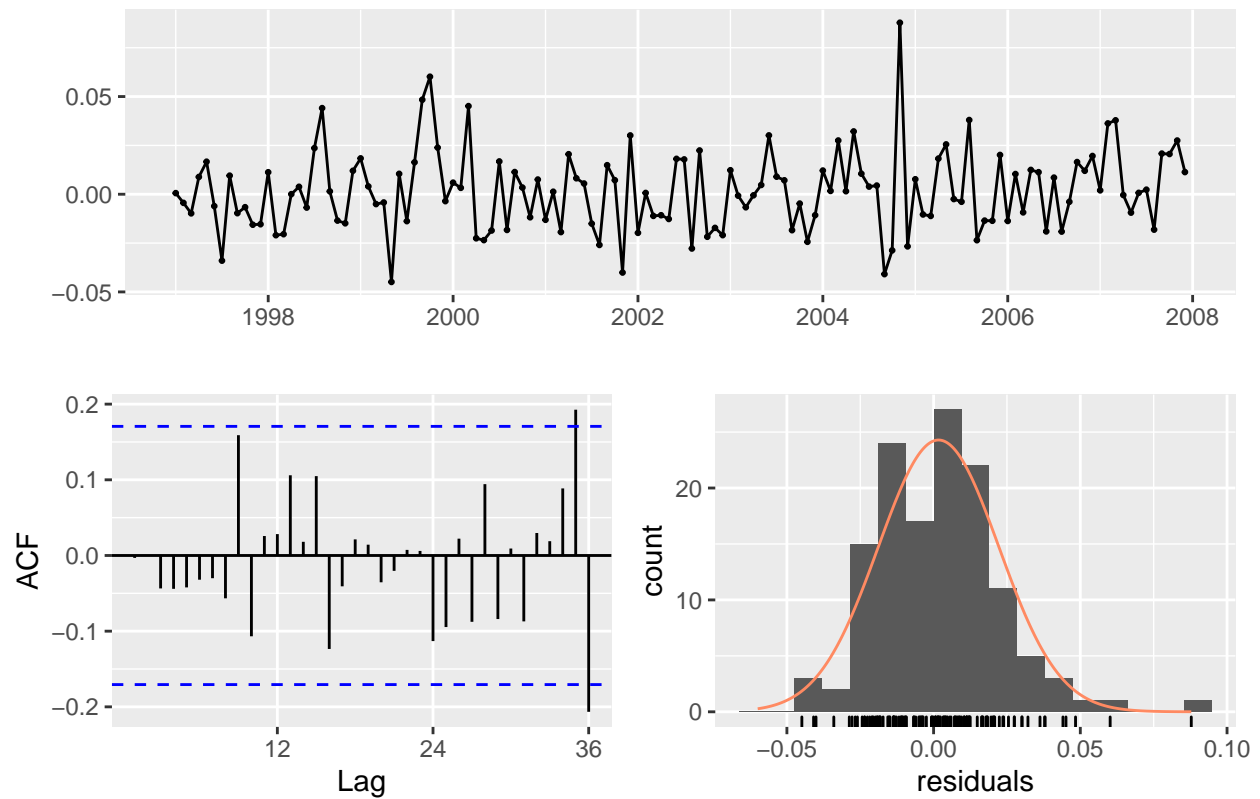
Residuals from ETS(M,A,A)



```
##
##  Ljung-Box test
##
## data:  Residuals from ETS(M,A,A)
## Q* = 12.469, df = 8, p-value = 0.1315
##
## Model df: 16.    Total lags used: 24
```

```
checkresiduals(fit5.ets)
```

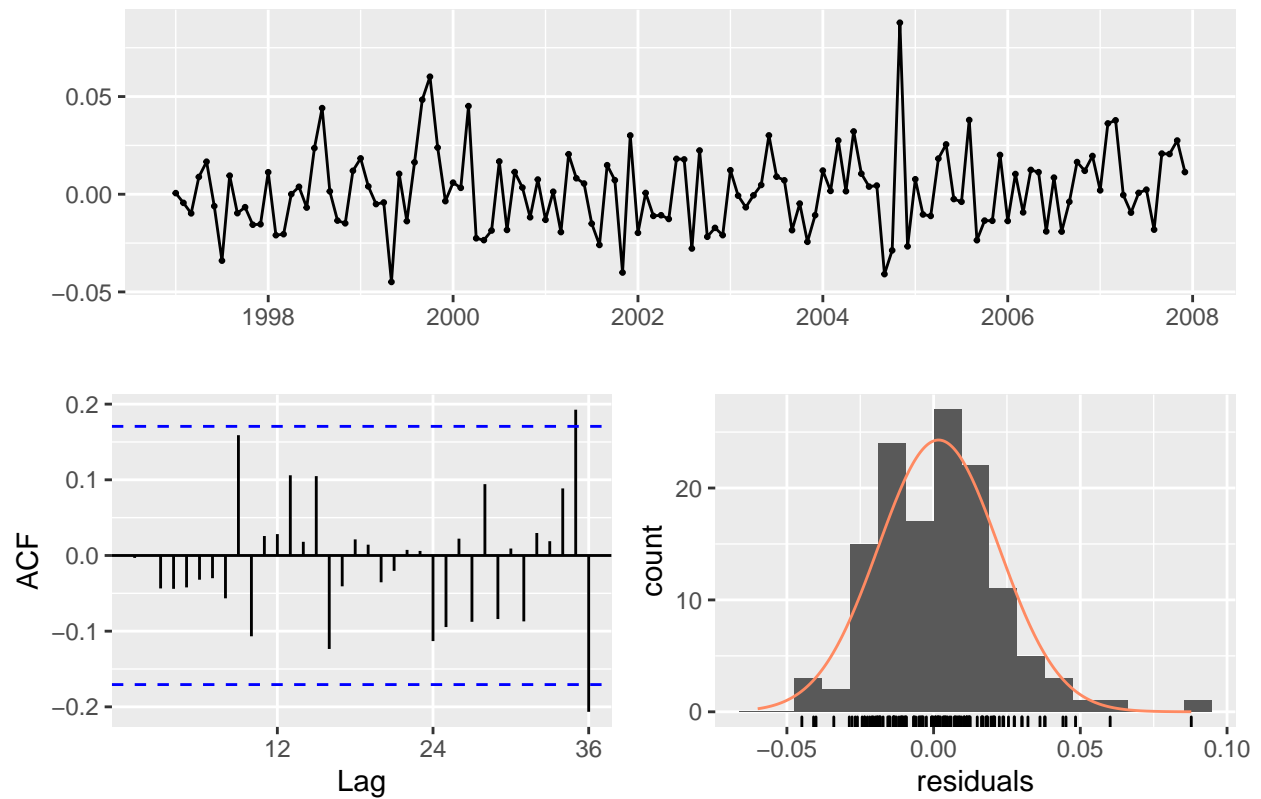
Residuals from ETS(M,Ad,M)



```
##
##  Ljung-Box test
##
## data:  Residuals from ETS(M,Ad,M)
## Q* = 15.447, df = 7, p-value = 0.03068
##
## Model df: 17.   Total lags used: 24
```

```
checkresiduals(fit6.ets)
```

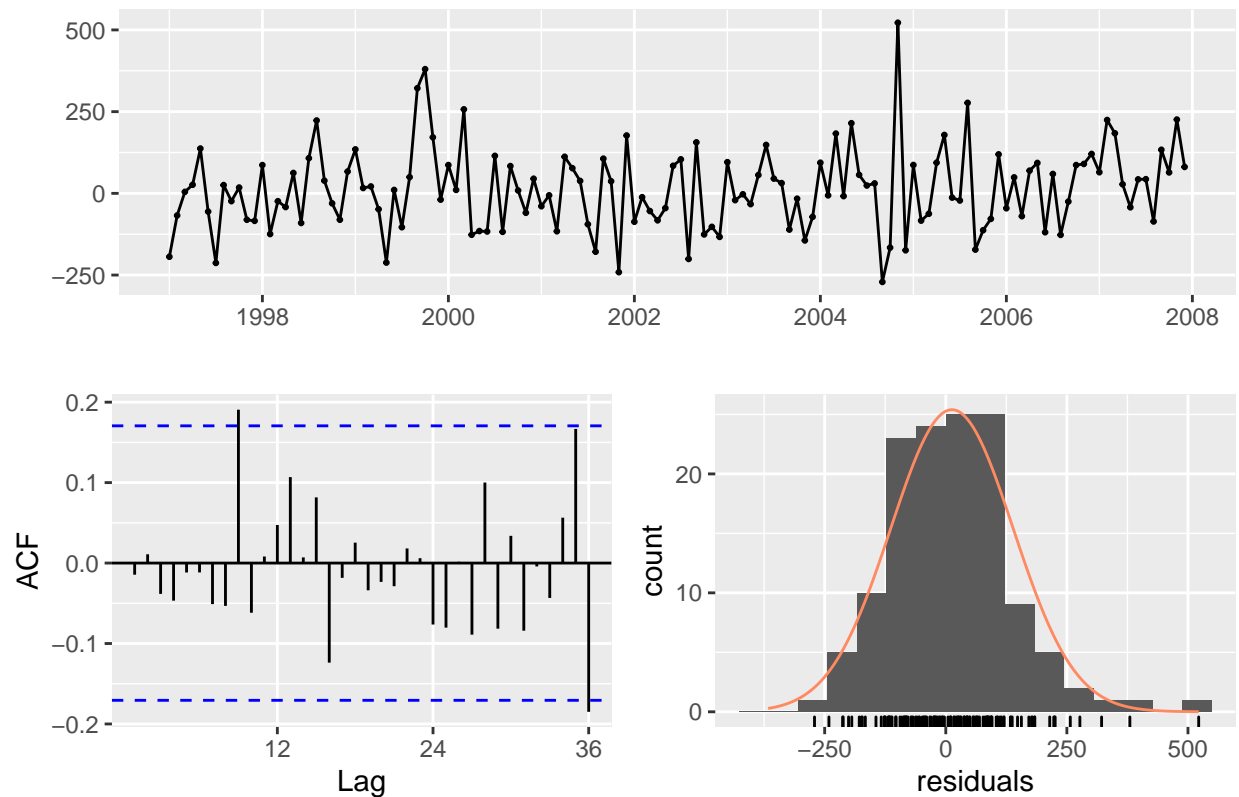
Residuals from ETS(M,Ad,M)



```
##
##  Ljung-Box test
##
## data:  Residuals from ETS(M,Ad,M)
## Q* = 15.447, df = 7, p-value = 0.03068
##
## Model df: 17.   Total lags used: 24
```

```
checkresiduals(fit7.ets)
```

Residuals from ETS(A,N,A)




```
##
##  Ljung-Box test
##
## data:  Residuals from ETS(A,N,A)
## Q* = 14.089, df = 10, p-value = 0.169
##
## Model df: 14.    Total lags used: 24
```

```
summary(fit4.ets)
```

```
## ETS(M,A,A)
##
## Call:
## ets(y = yvr.train, model = "MAA", damped = NULL)
##
## Smoothing parameters:
##   alpha = 0.8082
##   beta  = 1e-04
##   gamma = 1e-04
##
## Initial states:
##   l = 5803.689
##   b = 17.436
##   s = -37.9185 -34.4259 -248.9098 320.0811 666.5463 231.972
##       63.7197 -351.3296 -188.4109 -568.7838 53.8096 93.6498
```

```
##
##   sigma: 0.0218
##
##      AIC      AICc      BIC
## 1957.628 1962.996 2006.636
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -6.611667 126.3696 99.4423 -0.1458104 1.602955 0.3752775
##              ACF1
## Training set -0.02754951
```

```
# visualize the forecasts
plot(yvr.train, xlim=c(1997, 2011), ylim=c(5000, 8800), main='YVR Test data forecast',
     xlab='Years', ylab='Monthly Energy Use at YVR (1000s kWh)')
lines(yvr.test)
lines(fitted(fit4.ets), col='blue')
lines(forecast(fit4.ets, h=36)$mean, col='blue', lwd=2)
lines(forecast(fit4.ets, h=36)$upper[,2], col='blue', lwd=1, lty='dashed')
lines(forecast(fit4.ets, h=36)$lower[,2], col='blue', lwd=1, lty='dashed')
```



```
# accuracy scores
accuracy(forecast(fit4.ets, h=36), yvr.test)
```

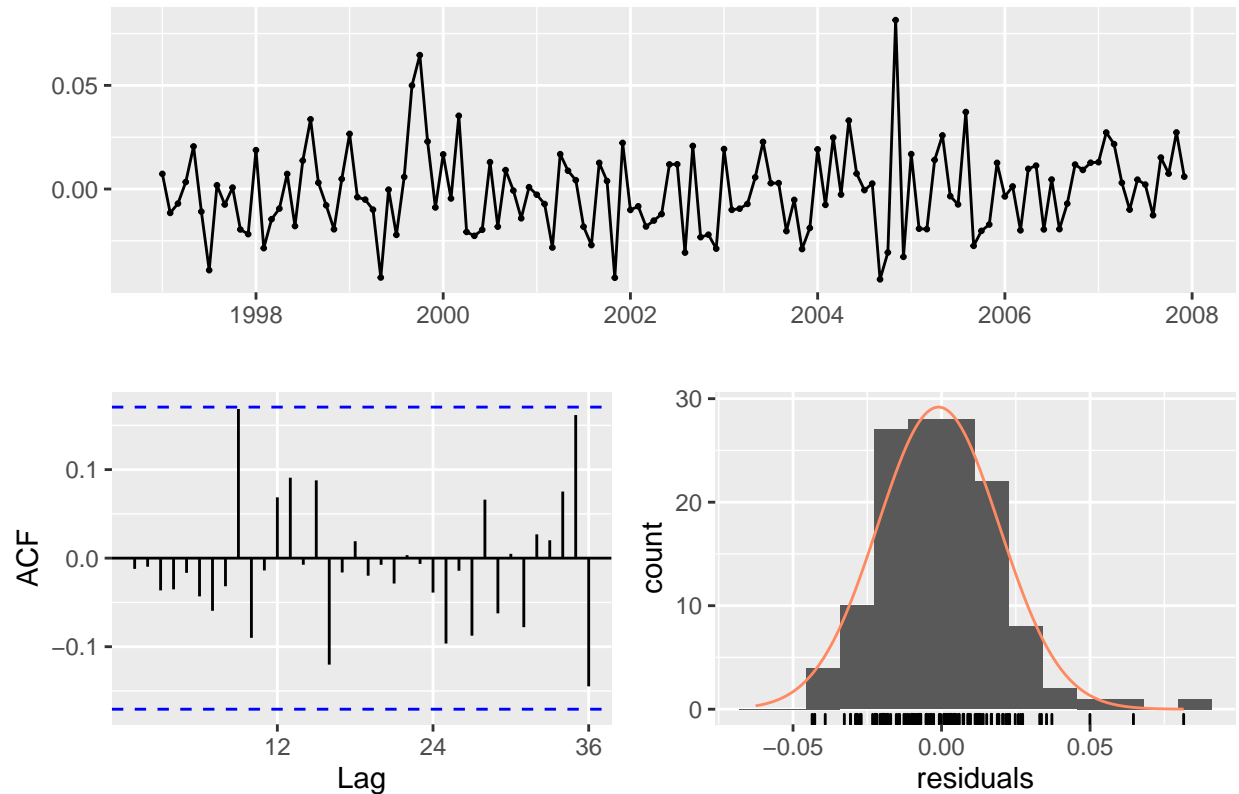
```
##              ME      RMSE      MAE      MPE      MAPE      MASE
```



```
## Training set -6.611667 126.3696 99.4423 -0.1458104 1.602955 0.3752775
## Test set    -20.414073 155.9826 123.1789 -0.3335670 1.614019 0.4648554
##              ACF1 Theil's U
## Training set -0.02754951      NA
## Test set      0.17004622 0.3285668
```

```
# residual diagnostics
checkresiduals(fit4.ets)
```

Residuals from ETS(M,A,A)



```
##
## Ljung-Box test
##
## data: Residuals from ETS(M,A,A)
## Q* = 12.469, df = 8, p-value = 0.1315
##
## Model df: 16. Total lags used: 24
```

```
# mean of residuals
mean(fit4.ets$residuals)
```

```
## [1] -0.001043112
```

```
# box ljung test of residual autocorrelations
Box.test(fit1.ets$residuals, type="Ljung", lag=24)
```

```
##
## Box-Ljung test
##
## data: fit1.ets$residuals
## X-squared = 13.308, df = 24, p-value = 0.9608
```

Additive Seasonality:

The seasonality seems quite constant across the time series, and did not amplify as the level increases.

Additive Trend without damping:

The trend does not look very exponential. Although the increasing energy use for the airport should become flatter (corresponding to damping), as the operation reaches limit (given the facility does not expand rapidly), we are uncertain should the capacity be reached in the next three years. Therefore we tried both damping and no damping. The version without damping has average residual closer to 0, so we chose no damping.

Multiplicative Errors:

There is no way to tell whether errors are additive or multiplicative from the original time series. Therefore, we tried both and Multiplicative Error has slightly less significant autocorrelations in the residuals.

The ETS model has forecasted the test data much more accurately than either of the basic methods, since all errors from ETS are much lower.

The MASE from the ETS model is less than 1, indicating the average test error is even lower than the training error.

The MAE for the ETS model shows that on average, the forecast on test data is about 123, 000 kilowatt hour off. 123 kwh per month is approximately 4 days of electricity consumption for a typical US family (refer to the bibliographical for the source). Such size of error for an airport can be negligible. Same can be inferred from RMSE.

In all, the ETS model of our choice did a good job on forecasting test data, which can transfer to good generalizability.

Zero MeanThe mean of time plot is approximately zero.

Constant VarianceOverall the variance is equal except for the 2 huge positive spike at around September 1999 just before year 2005. The spikes might be due to chance.

a) mean of residuals = -0.001043112

b)Histogram looks normal

The ACF plot looks perfect. We have no significant autocorrelations.

Box Ljung test for autocorrelations

H0: the first 24 autocorrelations are not significantly different from a white noise process

HA: the first 24 autocorrelations are significantly different from a white noise process

Test statistics = 13.308


p-value = 0.9608

Decision: We fail to reject null because p-value is bigger than 0.05

Conclusion: We conclude that the first 24 autocorrelations are not significantly different from a white noise process.

```
# differencing
yvr.train %>%
  diff(lag=12) %>%
  diff() %>%
  ggtsdisplay(main='Seasonally and Non-seasonally Differenced')
```

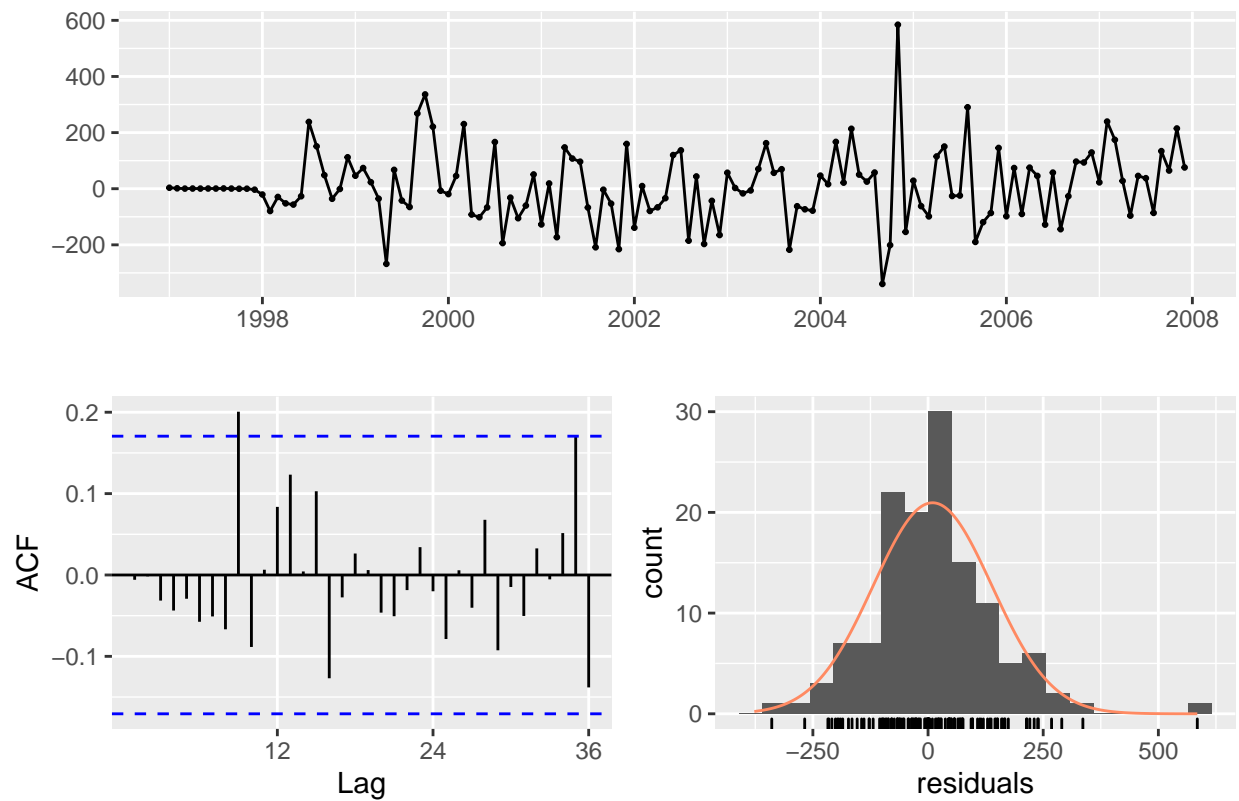
Seasonally and Non-seasonally Differenced



```
fit1.arma = Arima(yvr.train, order = c(1,1, 1), seasonal = c(0,1,1), include.constant = TRUE)
fit2.arma = Arima(yvr.train, order = c(2,1, 1), seasonal = c(1,1,1), include.constant = TRUE)
fit3.arma = Arima(yvr.train, order = c(0,1, 1), seasonal = c(1,1,1), include.constant = TRUE)
fit4.arma = Arima(yvr.train, order = c(2,1, 1), seasonal = c(0,1,1), include.constant = TRUE)
fit5.arma = auto.arima(yvr.train)
fit6.arma = Arima(yvr.train, order = c(0, 1, 1), seasonal = c(0,1,1), include.constant = TRUE)
fit7.arma = Arima(yvr.train, order = c(2, 1, 2), seasonal = c(1,1,1), include.constant = TRUE)
fit8.arma = Arima(yvr.train, order = c(0, 1, 0), seasonal = c(0,1,1), include.constant = TRUE)
fit9.arma = Arima(yvr.train, order = c(1,1, 0), seasonal = c(0,1,1), include.constant = TRUE)
```


```
checkresiduals(fit1.arma)
```

Residuals from ARIMA(1,1,1)(0,1,1)[12]




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(1,1,1)(0,1,1)[12]
## Q* = 17.561, df = 21, p-value = 0.6766
##
## Model df: 3.   Total lags used: 24
```

```
checkresiduals(fit2.arima)
```



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(2,1,1)(1,1,1)[12]
## Q* = 17.932, df = 19, p-value = 0.527
##
## Model df: 5.   Total lags used: 24
```

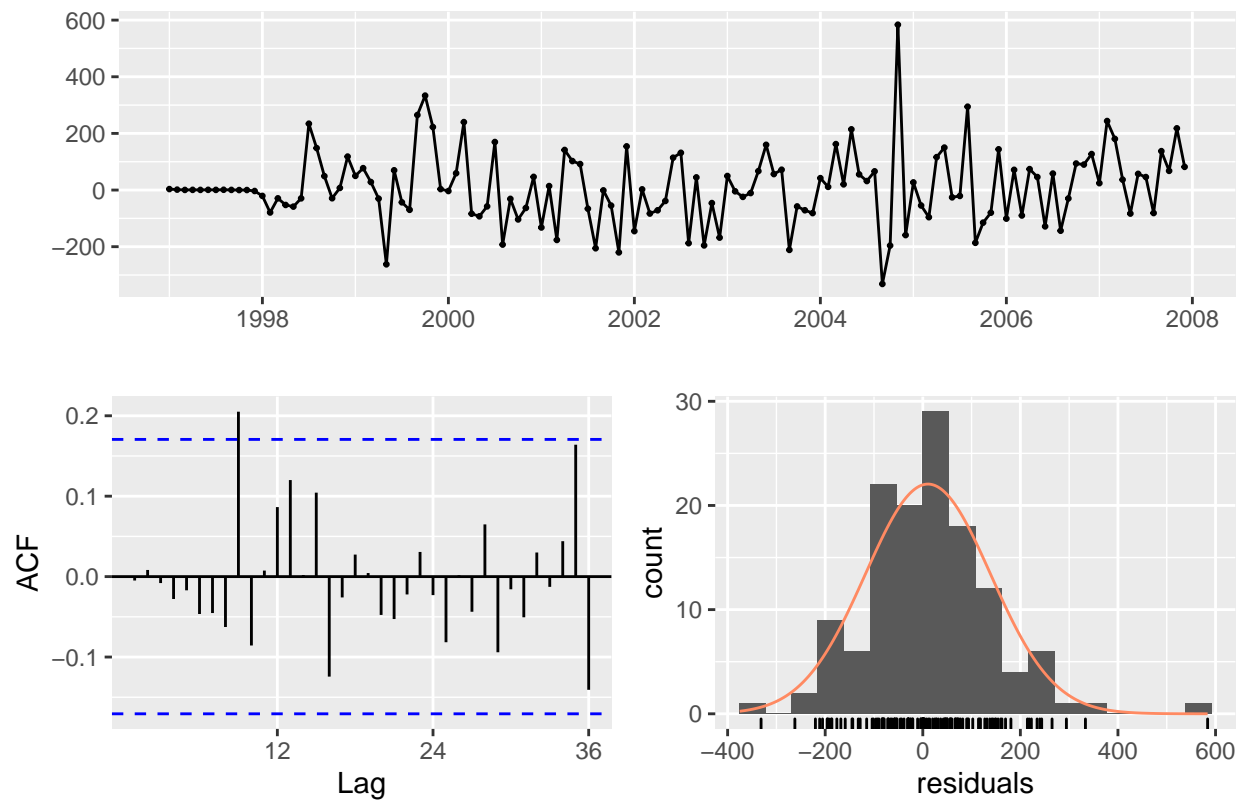
```
checkresiduals(fit3.arima)
```



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,1,1)(1,1,1)[12]
## Q* = 18.752, df = 21, p-value = 0.601
##
## Model df: 3.   Total lags used: 24
```

```
checkresiduals(fit4.arima)
```




Residuals from ARIMA(2,1,1)(0,1,1)[12]



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(2,1,1)(0,1,1)[12]
## Q* = 17.033, df = 20, p-value = 0.6508
##
## Model df: 4.   Total lags used: 24
```

```
checkresiduals(fit5.arima)
```

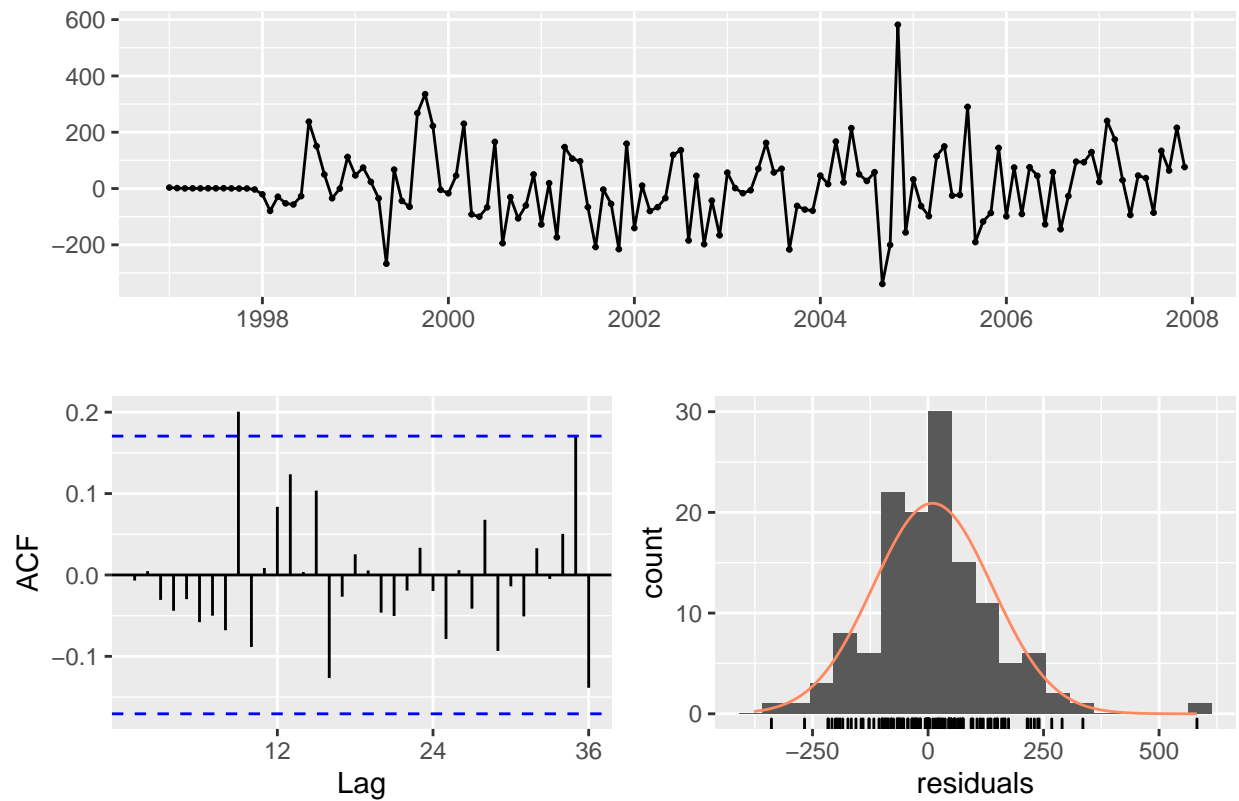
Residuals from ARIMA(1,1,0)(2,1,0)[12]



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(1,1,0)(2,1,0)[12]
## Q* = 22.208, df = 21, p-value = 0.3876
##
## Model df: 3.   Total lags used: 24
```

```
checkresiduals(fit6.arima)
```





Residuals from ARIMA(0,1,1)(0,1,1)[12]



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,1,1)(0,1,1)[12]
## Q* = 17.582, df = 22, p-value = 0.7305
##
## Model df: 2.   Total lags used: 24
```

```
checkresiduals(fit7.arima)
```




Residuals from ARIMA(2,1,2)(1,1,1)[12]



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(2,1,2)(1,1,1)[12]
## Q* = 16.895, df = 18, p-value = 0.5303
##
## Model df: 6.   Total lags used: 24
```

```
checkresiduals(fit8.arima)
```


Residuals from ARIMA(0,1,0)(0,1,1)[12]



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,1,0)(0,1,1)[12]
## Q* = 26.433, df = 23, p-value = 0.2808
##
## Model df: 1.   Total lags used: 24
```

```
checkresiduals(fit9.arima)
```

Residuals from ARIMA(1,1,0)(0,1,1)[12]




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(1,1,0)(0,1,1)[12]
## Q* = 17.722, df = 22, p-value = 0.7223
##
## Model df: 2.   Total lags used: 24
```

```
# check parameters
summary(fit9.arima)
```

```
## Series: yvr.train
## ARIMA(1,1,0)(0,1,1)[12]
##
## Coefficients:
##          ar1      sma1
##        -0.1972 -0.8927
## s.e.    0.0905  0.1794
##
## sigma^2 estimated as 18824:  log likelihood=-762.54
## AIC=1531.08  AICc=1531.29  BIC=1539.42
##
## Training set error measures:
##              ME    RMSE    MAE    MPE    MAPE    MASE
## Training set 9.495075 129.171 93.95688 0.1257726 1.494625 0.3545766
```

```
##                               ACF1
## Training set -0.01035696
```

```
# visualize the forecasts
plot(yvr.train, xlim=c(1997, 2011), ylim=c(5000, 8800), main='YVR Test data forecast',
     xlab='Years', ylab='Monthly Energy Use at YVR (1000s kWh)')
lines(yvr.test)
lines(fitted(fit9.arima), col='blue')
lines(forecast(fit9.arima, h=36)$mean, col='blue', lwd=2)
lines(forecast(fit9.arima, h=36)$upper[,2], col='blue', lwd=1, lty='dashed')
lines(forecast(fit9.arima, h=36)$lower[,2], col='blue', lwd=1, lty='dashed')
```

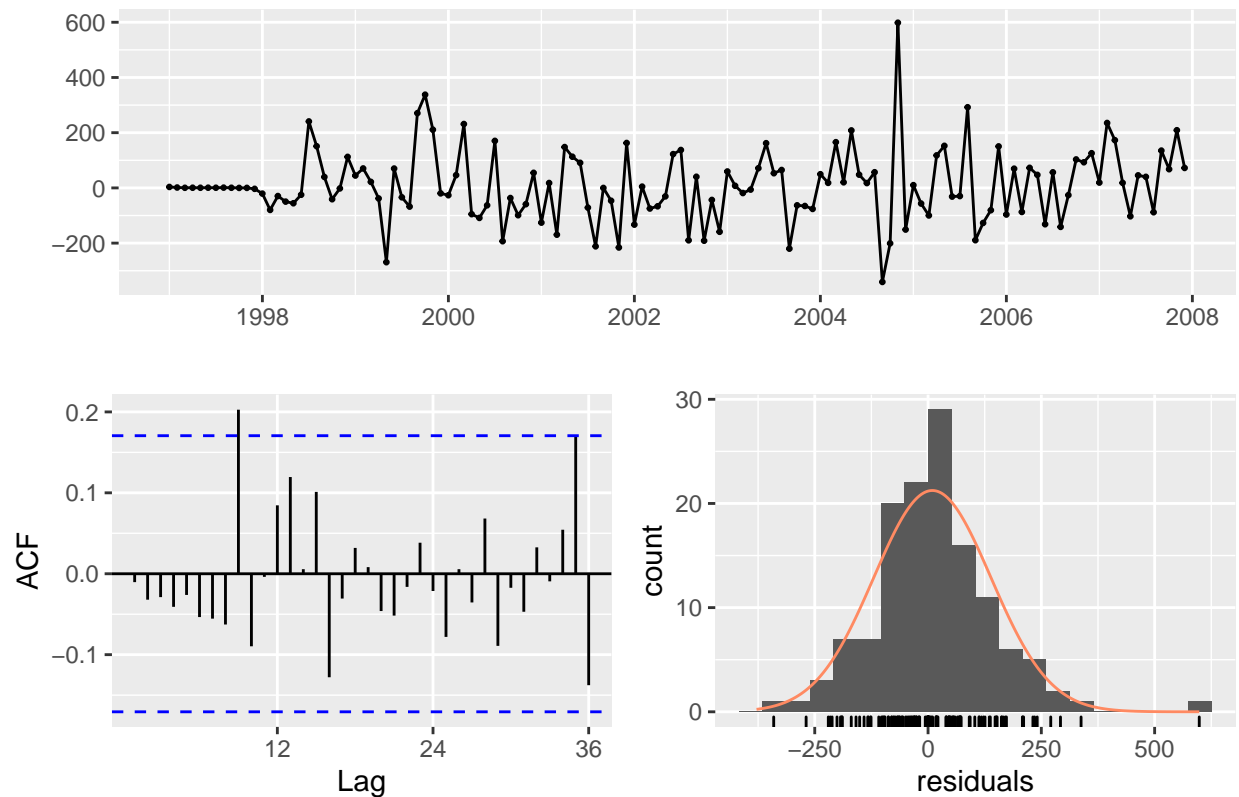


```
# accuracy measures
accuracy(forecast(fit9.arima, h=36), yvr.test)
```

```
##                               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  9.495075 129.1710  93.95688  0.1257726  1.494625  0.3545766
## Test set     -9.243577 159.9366 125.09305 -0.2000004  1.633898  0.4720790
##                               ACF1 Theil's U
## Training set -0.01035696      NA
## Test set     0.22421431 0.3373746
```

```
# residual diagnostics
checkresiduals(fit9.arima)
```

Residuals from ARIMA(1,1,0)(0,1,1)[12]



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(1,1,0)(0,1,1)[12]
## Q* = 17.722, df = 22, p-value = 0.7223
##
## Model df: 2.   Total lags used: 24
```

```
# mean of residuals
mean(fit9.arma$residuals)
```

```
## [1] 9.495075
```

```
# box ljung test of residual autocorrelations
Box.test(fit9.arma$residuals, type="Ljung", lag=24)
```

```
##
##  Box-Ljung test
##
## data:  fit9.arma$residuals
## X-squared = 17.722, df = 24, p-value = 0.8163
```

SAR =1:

The PACF looks like it has exponential decay at lag 12, 24, and 36 (though for lag 36 it might not be significant). There is a significant lag 12 for ACF.

AR=1:

The PACF has a significant lag 1. Though ACF also has a significant lag 1, the overall pattern is more complex. While suspecting it to be AR(1), I tried MA(1) as well. The residual plots of the two models are quite the same. Nevertheless, I still picked AR(1) as the ACF and PACF plots are more like it.

The ARIMA model has forecasted the test data much more accurately than either of the basic methods, since all errors from ARIMA are much lower. The MASE from the ARIMA model is less than 1, indicating the average test error is even lower than the training error. The MAE for the ARIMA model shows that on average, the forecast on test data is about 125,000 kilowatt hour off. 125 kwh per month is approximately 4 days of electricity consumption for a typical US family (refer to the bibliographical for the source). Such size of error for an airport can be negligible. Same can be inferred from RMSE. In all, the ARIMA model of our choice did a good job on forecasting test data, which can transfer to good generalizability.

- a) Mean of residuals = 9.495075. The slightly positive mean could be due to the positive spike of residual just before Year 2005.
- b) Residuals are a little positively skewed, mainly due to the existence of high residuals (of about 620, which could be outliers). Apart from the slight skewness, the histogram looks normal overall.

Overall no significant Autocorrelations from the ACF plot There is a significant spike at lag 9 which could be due to chance. Overall there is no significant autocorrelations.

Box Ljung test for autocorrelations

H0: the first 24 autocorrelations are not significantly different from a white noise process

HA: the first 24 autocorrelations are significantly different from a white noise process

Test statistics = 17.722

p-value = 0.8163

Decision: We fail to reject null because p-value is bigger than 0.05

Conclusion: We conclude that the first 24 autocorrelations are not significantly different from a white noise process.

Comparisons:

Goodness of fit

The ETS model has better RMSE of training data, while the ARIMA model has better MAE of training data. Generally the goodness of fit does not say anything about forecasting ability. Therefore, here we are only checking training errors but not using them for decision making.

Generalizability

Overall the two models have very close test errors. The ETS model of our choice has slightly better test errors (for all metrics) than the ARIMA model, while such differences are negligible to an airport.

Residual Diagnostics

First, the average residuals from ETS is -0.001043112, which is closer to 0 than the 9.495075 from ARIMA. Second, the residual distribution of the ETS model seems less skewed than that of the ARIMA model. Third, the autocorrelations of the residuals from ETS are less significant than those from the ARIMA (even if the significance spike in ARIMA might be due to chance). In all the residuals from ETS are closer to white noise than ARIMA, meaning the ETS model has captured more information than the ARIMA does.

Model Complexity

The two models have hyperparameters simple enough that run time and computational requirements will not be issues. Therefore, the two models have approximately the same level of complexity.

Conclusion

Overall I would choose the ETS model because it has lower test errors and more promising residuals.


```
# point forecasts and prediction intervals
options(digits = 2)
forecast(ets(yvr, model=fit4.ets, use.initial.values = TRUE), h=36)
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2011	8116	7883	8349	7760	8472
## Feb 2011	8094	7795	8392	7636	8551
## Mar 2011	7488	7146	7830	6965	8011
## Apr 2011	7886	7496	8276	7290	8483
## May 2011	7740	7312	8169	7085	8396
## Jun 2011	8173	7702	8644	7453	8893
## Jul 2011	8358	7848	8869	7578	9139
## Aug 2011	8810	8259	9362	7967	9654
## Sep 2011	8481	7897	9066	7587	9375
## Oct 2011	7930	7319	8540	6995	8864
## Nov 2011	8161	7521	8802	7182	9141
## Dec 2011	8175	7507	8843	7153	9197
## Jan 2012	8324	7628	9021	7260	9389
## Feb 2012	8302	7579	9024	7197	9406
## Mar 2012	7696	6954	8439	6561	8832
## Apr 2012	8094	7327	8861	6921	9268
## May 2012	7949	7160	8737	6742	9155
## Jun 2012	8381	7567	9195	7136	9626
## Jul 2012	8567	7728	9405	7284	9849
## Aug 2012	9019	8153	9884	7695	10343
## Sep 2012	8689	7801	9578	7331	10048
## Oct 2012	8138	7231	9045	6751	9525
## Nov 2012	8370	7441	9298	6950	9789
## Dec 2012	8383	7435	9332	6933	9834
## Jan 2013	8533	7563	9502	7050	10016
## Feb 2013	8510	7520	9500	6996	10024
## Mar 2013	7905	6899	8910	6367	9442
## Apr 2013	8302	7277	9327	6735	9870
## May 2013	8157	7115	9199	6563	9751
## Jun 2013	8589	7527	9651	6964	10214
## Jul 2013	8775	7692	9857	7119	10430
## Aug 2013	9227	8122	10331	7537	10916
## Sep 2013	8897	7774	10021	7180	10615
## Oct 2013	8346	7207	9485	6604	10088
## Nov 2013	8578	7421	9735	6808	10347
## Dec 2013	8592	7417	9766	6795	10388

```
plot(yvr,
      xlim=c(1997, 2013.8),
      ylim=c(5000, 11000),
      main='YVR new data forecast',
      xlab='Years', ylab='Monthly Energy Use at YVR (1000s kWh)')
```



```
lines(forecast(ets(yvr, model=fit4.ets, use.initial.values = TRUE), h=36)$mean, col='blue', lwd=2)
lines(forecast(ets(yvr, model=fit4.ets, use.initial.values = TRUE), h=36)$upper[,2], col='blue', lwd=1,
lines(forecast(ets(yvr, model=fit4.ets, use.initial.values = TRUE), h=36)$lower[,2], col='blue', lwd=1,
legend(1996.6, 11000, legend = c(' 3 year forecast', '95% prediction interval'), lty = c('solid', 'dash
```



Model Limitations

Limitation 1: Data LeakingThe first limitation regards to the model building process rather than the model itself. Back from the beginning of the project, we visualized the time plot of the entire data. Therefore we have already had a basic idea of the trend of the test data (upward), which would lead us to choose model with additive or multiplicative trend. Additionally, we did model selection based on test set, which overfits the test data. Such would increase the variance and therefore increases generalization error. (Our good performance on test data would not imply good performance on unseen future data)

Recommendation:Split data into training set, validation set, and test set. Use validation set for model selection to avoid overfitting test set.

Limitation 2: Only using past observations of the same variableThe caveat of the time series models is that they only use past observations of the same variable. Anything outside of seasonality, trend, and cycle cannot be explained by time series models.

Recommendation:Use exploratory models (details will be explained in the next question)

a) Idea 1:

Exploratory Model (Focus on Predictability)We can use all other variables available at hand (i.e. total area, passenger...) to predict month energy use. Since we focus on predictability, we should select high-performing black box models (eg: XGBoost, Random Forest...). Feature selections should be done to exclude

unimportant features. After that, to actually make forecast, we first build separate time series models to forecast the important explanatory variables you chose. Then we predict monthly energy use based on the forecasted explanatory variables. Such approach would explain patterns that are caused other than time

- b) Idea 2: Exploratory Model (Focus on Interpretability) Idea 2 has similar process to that in Idea 1. Should we focus on Interpretability of the model, we can only choose highly interpretable models (such as linear regression and decision trees). In the end we should be able to interpret coefficients on the linear regression or evaluate each split in a decision tree.
- c) Idea 3: Try ETS Models with multiplicative trend and additive seasonality Since ETS model forbids us to use combinations with additive seasonality and multiplicative trend, we need to first do seasonal adjustment so that the data has no